

---

# Supply Chain Dilemma

Gala Groceries is a technology-led grocery store chain based in the USA. They rely heavily on new technologies, such as IoT to give them a competitive edge over other grocery stores. Groceries are highly perishable items. If you overstock, you are wasting money on excessive storage and waste, but if you understock, then you risk losing customers.

This is logistic regression model built for Gala Groceries , a technology-led grocery store chain based in the USA to help them know and predict how to better stock grocery items that they sell

---

## Task 1 - Exploratory Data Analysis

---

### Section 1 - Importing Modules

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

---

### Section 2 - Data loading using Pandas

Loading `sample_sales_data.csv` dataset so that we can work with them in Python. For this notebook and all further notebooks, it will be assumed that the CSV files will be placed in the same file location as the notebook. If they are not, please adjust the directory within the `read_csv` method accordingly.

```
In [3]: sales_data = pd.read_csv(r"C:\Users\Mr.Hassan\DataspellProjects\Gala Foods\sample_sales_data.csv")
```

---

### Section 3 - Descriptive statistics

In this section, we try to gain a description of the data, that is: what columns are present, how many null values exist and what data types exists within each column.

To get started, this is an explanation of what the column names mean

- `transaction_id` = this is a unique ID that is assigned to each transaction
- `timestamp` = this is the datetime at which the transaction was made
- `product_id` = this is an ID that is assigned to the product that was sold. Each product has a unique ID
- `category` = this is the category that the product is contained within
- `customer_type` = this is the type of customer that made the transaction
- `unit_price` = the price that 1 unit of this item sells for
- `quantity` = the number of units sold for this product within this transaction
- `total` = the total amount payable by the customer
- `payment_type` = the payment method used by the customer

### Data Types

```
In [4]: sales_data.head()
```

---

Out[4]:

|   | Unnamed: 0 | transaction_id                       | timestamp           | product_id                           | category | customer_type | unit_price | quantity | total | payment_type |
|---|------------|--------------------------------------|---------------------|--------------------------------------|----------|---------------|------------|----------|-------|--------------|
| 0 | 0          | a1c82654-c52c-45b3-8ce8-4c2a1efe63ed | 2022-03-02 09:51:38 | 3bc6c1ea-0198-46de-9ffd-514ae3338713 | fruit    | gold          | 3.99       | 2        | 7.98  | e-wallet     |
| 1 | 1          | 931ad550-09e8-4da6-beaa-8c9d17be9c60 | 2022-03-06 10:33:59 | ad81b46c-bf38-41cf-9b54-5fe7f5eba93e | fruit    | standard      | 3.99       | 1        | 3.99  | e-wallet     |
| 2 | 2          | ae133534-6f61-4cd6-b6b8-d1c1d8d90aea | 2022-03-04 17:20:21 | 7c55cbd4-f306-4c04-a030-628cbe7867c1 | fruit    | premium       | 0.19       | 2        | 0.38  | e-wallet     |
| 3 | 3          | 157cebd9-aaf0-475d-8a11-7c8e0f5b76e4 | 2022-03-02 17:23:58 | 80da8348-1707-403f-8be7-9e6decccc883 | fruit    | gold          | 0.19       | 4        | 0.76  | e-wallet     |
| 4 | 4          | a81a6cd3-5e0c-44a2-826c-aea43e46c514 | 2022-03-05 14:32:43 | 7f5e86e6-f06f-45f6-bf44-27b095c9ad1d | fruit    | basic         | 4.49       | 2        | 8.98  | debit card   |

In [5]:

```
sales_data.info()
#There are no null values in this data set.

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7829 entries, 0 to 7828
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            7829 non-null   int64
1   transaction_id        7829 non-null   object
2   timestamp             7829 non-null   object
3   product_id           7829 non-null   object
4   category             7829 non-null   object
5   customer_type        7829 non-null   object
6   unit_price           7829 non-null   float64
7   quantity             7829 non-null   int64
8   total               7829 non-null   float64
9   payment_type         7829 non-null   object
dtypes: float64(2), int64(2), object(6)
memory usage: 611.8+ KB
```

Statistics

In [6]:

```
sales_data.describe()
#We might have skewed data columns here
```

Out[6]:

|       | Unnamed: 0  | unit_price  | quantity    | total       |
|-------|-------------|-------------|-------------|-------------|
| count | 7829.000000 | 7829.000000 | 7829.000000 | 7829.000000 |
| mean  | 3914.000000 | 7.819480    | 2.501597    | 19.709905   |
| std   | 2260.181962 | 5.388088    | 1.122722    | 17.446680   |
| min   | 0.000000    | 0.190000    | 1.000000    | 0.190000    |
| 25%   | 1957.000000 | 3.990000    | 1.000000    | 6.570000    |
| 50%   | 3914.000000 | 7.190000    | 3.000000    | 14.970000   |
| 75%   | 5871.000000 | 11.190000   | 4.000000    | 28.470000   |
| max   | 7828.000000 | 23.990000   | 4.000000    | 95.960000   |

Section 4 - Visualisation

Now that we have computed some descriptive statistics of the dataset, let's create some visualisations.

In [7]:

```
def plot_continuous_distribution(data: pd.DataFrame = None, column: str = None, height: int = 8):
    sns.displot(data, x=column, kde=True, height=height, aspect=height/5).set(title=f'Distribution of {column}')


def plot_categorical_distribution(data: pd.DataFrame = None, column: str = None, height: int = 8, aspect: int = 5):
    sns.catplot(data=data, x=column, kind='count', height=height, aspect=aspect, order=data[column].value_counts)

def correlation_plot(data: pd.DataFrame = None):
    corr = data.corr()
    corr.style.background_gradient(cmap='coolwarm')
    sns.heatmap(corr, xticklabels=corr.columns.values, yticklabels=corr.columns.values, annot = True, annot_kws:
    # Axis ticks size
    plt.xticks(fontsize=10)
    plt.yticks(fontsize=10)
    plt.show()
```


To analyse the dataset, below are snippets of helper functions to visualise different columns within the dataset.

- plot\_continuous\_distribution = this is to visualise the distribution of numeric columns

```
In [8]: plot_continuous_distribution(sales_data,"unit_price")
```




```
In [9]: plot_continuous_distribution(sales_data,"total")
```




- plot\_categorical\_distribution = this is to visualise the distribution of categorical columns


```
In [10]: plot_categorical_distribution(sales_data,"category")
```



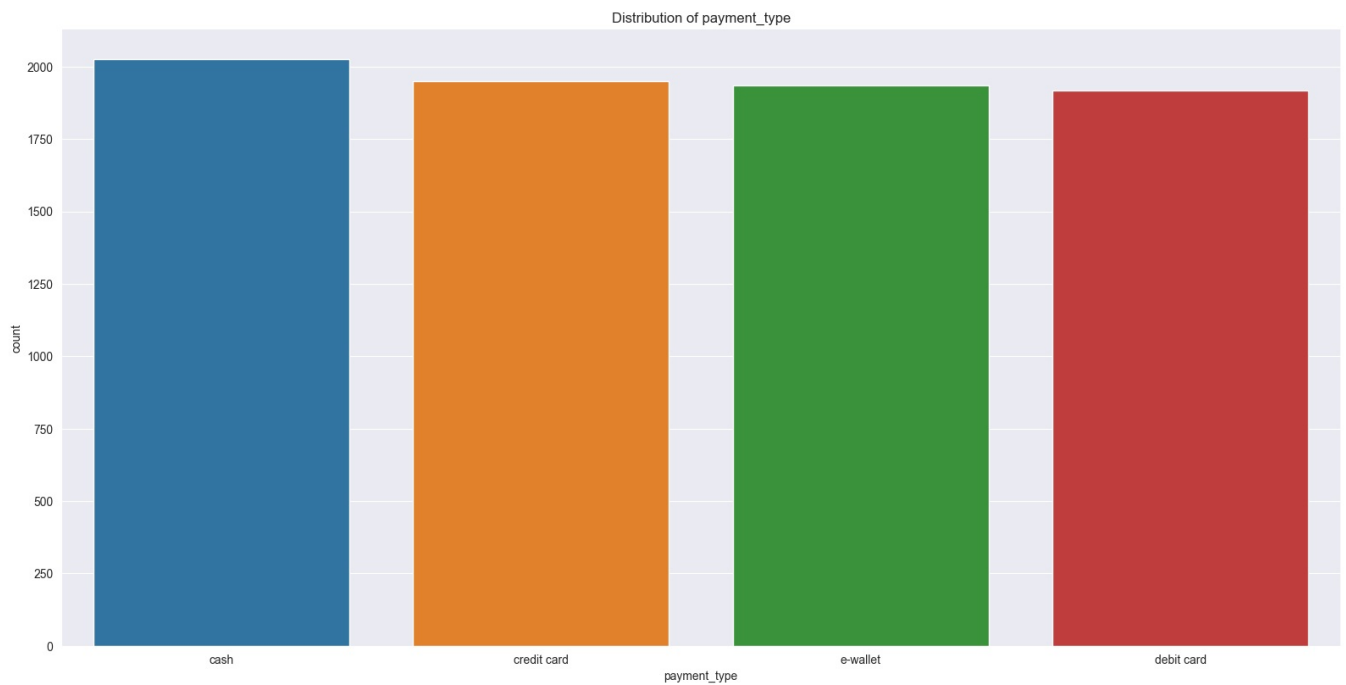
```
In [11]: plot_categorical_distribution(sales_data,"customer_type")
```



```
In [12]: plot_categorical_distribution(sales_data,"quantity")
```




```
In [13]: plot_categorical_distribution(sales_data,"payment_type")
```

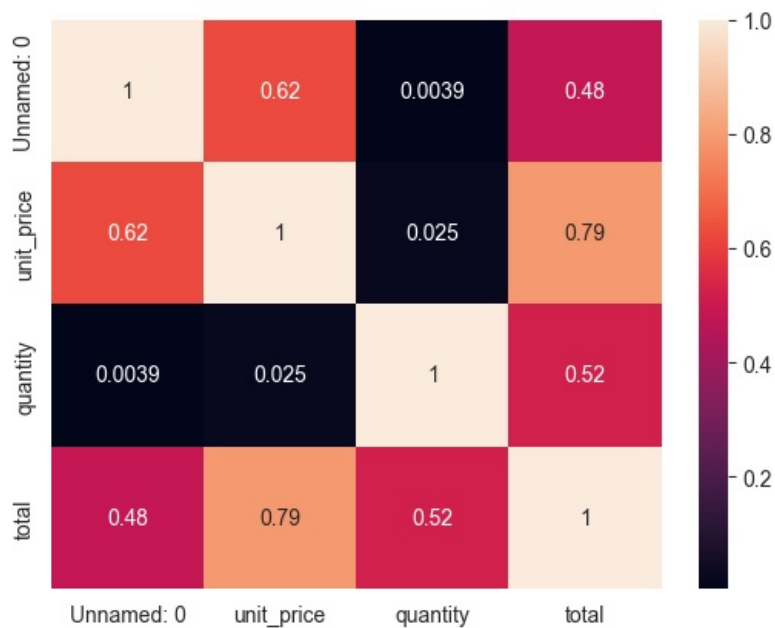


- correlation\_plot = this is to plot the correlations between the numeric columns within the data

```
In [14]: plot_continuous_distribution(sales_data,"total")
```



```
In [15]: correlation_plot(sales_data)
```



## Section 5 - Summary

We have completed an initial exploratory data analysis on the sample of data provided. We should now have a solid understanding of the data. I found the following insights as part of the analysis:

- Fruit & vegetables are the 2 most frequently bought product categories
- Non-members are the most frequent buyers within the store
- Cash is the most frequently used payment method

### DISCREPANCY

The client wants to know "How to better stock the items that they sell". From this dataset, it is impossible to answer that question. In order to make the next step on this project with the client, it is clear that:

- We need more rows of data. The current sample is only from 1 store and 1 week worth of data
- We need to frame the specific problem statement that we want to solve. The current business problem is too broad, we should

narrow down the focus in order to deliver a valuable end product

- We need more features. Based on the problem statement that we move forward with, we need more columns (features) that may help us to understand the outcome that we're solving for

Cont. from Task 1(Exploratory Data Analysis)

## Task 2 - Modeling

### Section 6 - Importing Modules to be used

```
In [16]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

### Section 7 - Data loading

Similar to before, let's load our data from the 3 datasets provided Loading

`sample_sales_data.csv` , `sensor_stock_levels` , `sensor_storage_temperature` dataset so that we can work with them in Python. For this notebook and all further notebooks, it will be assumed that the CSV files will be placed in the same file location as the notebook. If they are not, please adjust the directory within the `read_csv` method accordingly.

```
In [17]: sales_data = pd.read_csv(r"C:\Users\Mr.Hassan\DataSpellProjects\Gala Foods\sales.csv")
sensor_storage_temperature = pd.read_csv(r"C:\Users\Mr.Hassan\DataSpellProjects\Gala Foods\sensor_storage_tempe
sensor_stock_levels = pd.read_csv(r"C:\Users\Mr.Hassan\DataSpellProjects\Gala Foods\sensor_stock_levels.csv")
```

```
In [18]: sales_data.drop(columns=["Unnamed: 0"],inplace=True)
sales_data
```

|      | transaction_id                       | timestamp           | product_id                           | category          | customer_type | unit_price | quantity | total | payment_type |
|------|--------------------------------------|---------------------|--------------------------------------|-------------------|---------------|------------|----------|-------|--------------|
| 0    | a1c82654-c52c-45b3-8ce8-4c2a1efe63ed | 2022-03-02 09:51:38 | 3bc6c1ea-0198-46de-9ffd-514ae3338713 | fruit             | gold          | 3.99       | 2        | 7.98  | e-wallet     |
| 1    | 931ad550-09e8-4da6-beaa-8c9d17be9c60 | 2022-03-06 10:33:59 | ad81b46c-bf38-41cf-9b54-5fe7f5eba93e | fruit             | standard      | 3.99       | 1        | 3.99  | e-wallet     |
| 2    | ae133534-6f61-4cd6-b6b8-d1c1d8d90aea | 2022-03-04 17:20:21 | 7c55cbd4-f306-4c04-a030-628cbe7867c1 | fruit             | premium       | 0.19       | 2        | 0.38  | e-wallet     |
| 3    | 157cebd9-aaf0-475d-8a11-7c8e0f5b76e4 | 2022-03-02 17:23:58 | 80da8348-1707-403f-8be7-9e6deecce883 | fruit             | gold          | 0.19       | 4        | 0.76  | e-wallet     |
| 4    | a81a6cd3-5e0c-44a2-826c-aea43e46c514 | 2022-03-05 14:32:43 | 7f5e86e6-f06f-45f6-bf44-27b095c9ad1d | fruit             | basic         | 4.49       | 2        | 8.98  | debit card   |
| ...  | ...                                  | ...                 | ...                                  | ...               | ...           | ...        | ...      | ...   | ...          |
| 7824 | 6c19b9fc-f86d-4526-9dfe-d8027a4d13ee | 2022-03-03 18:22:09 | bc6187a9-d508-482b-9ca6-590d1cc7524f | cleaning products | basic         | 14.19      | 2        | 28.38 | e-wallet     |
| 7825 | 1c69824b-e399-4b79-a5e7-04a3a7db0681 | 2022-03-04 19:14:46 | 707e4237-191c-4cc9-85af-383a6c1cb2ab | cleaning products | standard      | 16.99      | 1        | 16.99 | credit card  |
| 7826 | 79aee7d6-1405-4345-9a15-92541e9e1e74 | 2022-03-03 14:00:09 | a9325c1a-2715-41df-b7f4-3078fa5ecd97 | cleaning products | basic         | 14.19      | 2        | 28.38 | credit card  |
| 7827 | e5cc4f88-e5b7-4ad5-bc1b-12a828a14f55 | 2022-03-04 15:11:38 | 707e4237-191c-4cc9-85af-383a6c1cb2ab | cleaning products | basic         | 16.99      | 4        | 67.96 | cash         |
| 7828 | afd70b4f-ee21-402d-8d8f-0d9e13c2bea6 | 2022-03-06 13:50:36 | d6ccd088-11be-4c25-aa1f-ea87c01a04db | cleaning products | non-member    | 14.99      | 4        | 59.96 | debit card   |

7829 rows × 9 columns

```
In [19]: sensor_stock_levels.drop(columns=["Unnamed: 0"],inplace=True)
sensor_stock_levels
```

Out[19]:

|       |                                      | id                  | timestamp                            | product_id | estimated_stock_pct |
|-------|--------------------------------------|---------------------|--------------------------------------|------------|---------------------|
| 0     | 4220e505-c247-478d-9831-6b9f87a4488a | 2022-03-07 12:13:02 | f658605e-75f3-4fed-a655-c0903f344427 | 0.75       |                     |
| 1     | f2612b26-fc82-49ea-8940-0751fdd4d9ef | 2022-03-07 16:39:46 | de06083a-f5c0-451d-b2f4-9ab88b52609d | 0.48       |                     |
| 2     | 989a287f-67e6-4478-aa49-c3a35dac0e2e | 2022-03-01 18:17:43 | ce8f3a04-d1a4-43b1-a7c2-fa1b8e7674c8 | 0.58       |                     |
| 3     | af8e5683-d247-46ac-9909-1a77bdebefb2 | 2022-03-02 14:29:09 | c21e3ba9-92a3-4745-92c2-6faef73223f7 | 0.79       |                     |
| 4     | 08a32247-3f44-4002-85fb-c198434dd4bb | 2022-03-02 13:46:18 | 7f478817-aa5b-44e9-9059-8045228c9eb0 | 0.22       |                     |
| ...   | ...                                  | ...                 | ...                                  | ...        |                     |
| 14995 | b9bf6788-09f3-490b-959b-dc5b55edb4b6 | 2022-03-04 10:52:50 | e37658de-3649-4ddb-9c73-b868dd69d3fe | 0.66       |                     |
| 14996 | 9ff1cc01-020f-491a-bafd-13552dccff44 | 2022-03-02 12:25:48 | fbeb39cc-8cd0-4143-bdfb-77658a02dec9 | 0.99       |                     |
| 14997 | 4d8101de-e8a2-4af9-9764-7a3a22aa7084 | 2022-03-03 17:36:44 | 8e21dcec-d775-4969-8334-05a37a5fd189 | 0.72       |                     |
| 14998 | 5f2a7b1e-b3c4-4395-8425-c960e22f701d | 2022-03-02 19:42:47 | 9708cf5b-aa69-4320-a013-9d234c40e63f | 0.95       |                     |
| 14999 | af6f4493-e49d-4dcb-951d-308e6cce267b | 2022-03-06 17:18:27 | 3bc6c1ea-0198-46de-9ffd-514ae3338713 | 0.75       |                     |

15000 rows × 4 columns

In [20]:

```
sensor_storage_temperature.drop(columns=["Unnamed: 0"],inplace=True)
sensor_storage_temperature
```

Out[20]:

|       |                                      | id                  | timestamp | temperature |
|-------|--------------------------------------|---------------------|-----------|-------------|
| 0     | d1ca1ef8-0eac-42fc-af80-97106efc7b13 | 2022-03-07 15:55:20 | 2.96      |             |
| 1     | 4b8a66c4-0f3a-4f16-826f-8cf9397e9d18 | 2022-03-01 09:18:22 | 1.88      |             |
| 2     | 3d47a0c7-1e72-4512-812f-b6b5d8428cf3 | 2022-03-04 15:12:26 | 1.78      |             |
| 3     | 9500357b-ce15-424a-837a-7677b386f471 | 2022-03-02 12:30:42 | 2.18      |             |
| 4     | c4b61fec-99c2-4c6d-8e5d-4edd8c9632fa | 2022-03-05 09:09:33 | 1.38      |             |
| ...   |                                      | ...                 | ...       | ...         |
| 23885 | 17bcff56-9965-4e9f-ad5f-107f0f3be93f | 2022-03-01 10:40:43 | -1.46     |             |
| 23886 | 51d4eb44-04bd-4d6a-b777-0653bc173303 | 2022-03-05 17:07:49 | -19.37    |             |
| 23887 | bbcacfc4-3b59-47ee-b9e1-7dd3bd588748 | 2022-03-01 16:15:41 | -2.89     |             |
| 23888 | 5c4d567b-4bcf-4fcd-86b7-e2db5de6e439 | 2022-03-07 14:44:52 | -2.56     |             |
| 23889 | 589c28e1-f1f3-4efb-af6d-9f194c4d7d5b | 2022-03-01 16:33:41 | 0.13      |             |

23890 rows × 3 columns

## Section 8 - Descriptive Statistics

In this section, we try to gain a description of the data, that is: what columns are present, how many null values exist and what data types exists within each column.

- SALES DATA

In [21]:

```
sales_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7829 entries, 0 to 7828
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   transaction_id   7829 non-null   object
1   timestamp        7829 non-null   object
2   product_id       7829 non-null   object
3   category         7829 non-null   object
4   customer_type    7829 non-null   object
5   unit_price       7829 non-null   float64
6   quantity         7829 non-null   int64
7   total            7829 non-null   float64
8   payment_type     7829 non-null   object
dtypes: float64(2), int64(1), object(6)
memory usage: 550.6+ KB
```

In [22]:

```
sales_data.describe()
```



Out[22]:

|       | unit_price  | quantity    | total       |
|-------|-------------|-------------|-------------|
| count | 7829.000000 | 7829.000000 | 7829.000000 |
| mean  | 7.819480    | 2.501597    | 19.709905   |
| std   | 5.388088    | 1.122722    | 17.446680   |
| min   | 0.190000    | 1.000000    | 0.190000    |
| 25%   | 3.990000    | 1.000000    | 6.570000    |
| 50%   | 7.190000    | 3.000000    | 14.970000   |
| 75%   | 11.190000   | 4.000000    | 28.470000   |
| max   | 23.990000   | 4.000000    | 95.960000   |

- sensor stock data

```
In [23]: sensor_stock_levels.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    15000 non-null  object
1   timestamp             15000 non-null  object
2   product_id            15000 non-null  object
3   estimated_stock_pct    15000 non-null  float64
dtypes: float64(1), object(3)
memory usage: 468.9+ KB
```

```
In [24]: sensor_stock_levels.describe()
```

Out[24]:

|       | estimated_stock_pct |
|-------|---------------------|
| count | 15000.000000        |
| mean  | 0.502735            |
| std   | 0.286842            |
| min   | 0.010000            |
| 25%   | 0.260000            |
| 50%   | 0.500000            |
| 75%   | 0.750000            |
| max   | 1.000000            |

- sensor storage temperature

```
In [25]: sensor_storage_temperature.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23890 entries, 0 to 23889
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    23890 non-null  object
1   timestamp             23890 non-null  object
2   temperature            23890 non-null  float64
dtypes: float64(1), object(2)
memory usage: 560.0+ KB
```

```
In [26]: sensor_storage_temperature.describe()
```

Out[26]:

|       | temperature  |
|-------|--------------|
| count | 23890.000000 |
| mean  | -0.207075    |
| std   | 11.217649    |
| min   | -30.990000   |
| 25%   | -2.860000    |
| 50%   | -1.000000    |
| 75%   | 1.840000     |
| max   | 34.990000    |

If we revisit the problem statement: Can we accurately predict the stock levels of products, based on sales

data and sensor data, on an hourly basis in order to more intelligently procure products from our suppliers.

The client indicates that they want the model to predict on an hourly basis. Looking at the data model, we can see that only column that we can use to merge the 3 datasets together is timestamp.

1. we must first transform the timestamp format to datetime format,
2. And then convert these timestamp columns in all 3 datasets to be based on the hour of the day.

```
In [27]: #A function that converts timestamp to datetime
def convert_to_datetime(data:pd.DataFrame=None, column:str=None):
    data[column] = pd.to_datetime(data[column],format="%Y-%m-%d %H:%M:%S")
```

```
In [28]: convert_to_datetime(sales_data,"timestamp")
convert_to_datetime(sensor_storage_temperature,"timestamp")
convert_to_datetime(sensor_stock_levels,"timestamp")
```

```
In [29]: #All timestamps have now been converted into datetimes
sensor_stock_levels.info()
sales_data.info()
sensor_storage_temperature.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    15000 non-null  object
1   timestamp             15000 non-null  datetime64[ns]
2   product_id            15000 non-null  object
3   estimated_stock_pct    15000 non-null  float64
dtypes: datetime64[ns](1), float64(1), object(2)
memory usage: 468.9+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7829 entries, 0 to 7828
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   transaction_id        7829 non-null  object
1   timestamp             7829 non-null  datetime64[ns]
2   product_id            7829 non-null  object
3   category              7829 non-null  object
4   customer_type         7829 non-null  object
5   unit_price            7829 non-null  float64
6   quantity              7829 non-null  int64
7   total                 7829 non-null  float64
8   payment_type          7829 non-null  object
dtypes: datetime64[ns](1), float64(2), int64(1), object(5)
memory usage: 550.6+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23890 entries, 0 to 23889
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    23890 non-null  object
1   timestamp             23890 non-null  datetime64[ns]
2   temperature           23890 non-null  float64
dtypes: datetime64[ns](1), float64(1), object(1)
memory usage: 560.0+ KB
```

```
In [30]: #We now convert these datetimes to hours
from datetime import datetime

def convert_timestamp_to_hourly(data: pd.DataFrame = None, column: str = None):
    new_ts = data[column].tolist()
    new_ts = map(lambda i:i.strftime('%Y-%m-%d %H:00:00'),new_ts)
    new_ts = [datetime.strptime(i, '%Y-%m-%d %H:00:00') for i in new_ts]
    data[column] = new_ts
```

```
In [31]: convert_timestamp_to_hourly(sales_data,"timestamp")
convert_timestamp_to_hourly(sensor_storage_temperature,"timestamp")
convert_timestamp_to_hourly(sensor_stock_levels,"timestamp")
```

Now all the timestamp columns have their minutes and seconds reduced to 00. The next thing to do, is to aggregate the datasets in order to combine rows which have the same value for timestamp .

For the sales data, we want to group the data by timestamp but also by product\_id. When we aggregate, we must choose which columns to aggregate by the grouping. For now, let's aggregate quantity.

```
In [32]: sales_agg = sales_data.groupby(['timestamp', 'product_id']).agg({"quantity": 'mean'}).reset_index()
```

```
sales_agg
```

```
Out[32]:
```

|      | timestamp           | product_id                           | quantity |
|------|---------------------|--------------------------------------|----------|
| 0    | 2022-03-01 09:00:00 | 00e120bb-89d6-4df5-bc48-a051148e3d03 | 3.00     |
| 1    | 2022-03-01 09:00:00 | 01f3cdd9-8e9e-4dff-9b5c-69698a0388d0 | 3.00     |
| 2    | 2022-03-01 09:00:00 | 03a2557a-aa12-4add-a6d4-77dc36342067 | 3.00     |
| 3    | 2022-03-01 09:00:00 | 049b2171-0eeb-4a3e-bf98-0c290c7821da | 3.50     |
| 4    | 2022-03-01 09:00:00 | 04da844d-8dba-4470-9119-e534d52a03a0 | 2.75     |
| ...  | ...                 | ...                                  | ...      |
| 6212 | 2022-03-07 19:00:00 | edf4ac93-4e14-4a3d-8c60-e715210cf3f9 | 3.00     |
| 6213 | 2022-03-07 19:00:00 | f01b189c-6345-4639-a8d1-89e1fc67c443 | 3.00     |
| 6214 | 2022-03-07 19:00:00 | f3bec808-bee0-4597-a129-53a3a2805a43 | 2.00     |
| 6215 | 2022-03-07 19:00:00 | fd66ac0b-3498-4613-8ec0-764686b0d864 | 1.00     |
| 6216 | 2022-03-07 19:00:00 | fd77b5cb-498c-40ca-95d1-0f87f13dd0d8 | 1.00     |

6217 rows × 3 columns

- We now have an aggregated sales data where each row represents a unique combination of hour during which the sales took place from that weeks worth of data and the product\_id.
- We summed the quantity and we took the mean average of the unit\_price.

For the stock data, we want to group it in the same way and aggregate the estimated\_stock\_pct.

```
In [33]: sensor_stock_levels_agg = sensor_stock_levels.groupby(['timestamp', 'product_id']).agg({'estimated_stock_pct':  
sensor_stock_levels_agg.head()
```

```
Out[33]:
```

|   | timestamp           | product_id                           | estimated_stock_pct |
|---|---------------------|--------------------------------------|---------------------|
| 0 | 2022-03-01 09:00:00 | 00e120bb-89d6-4df5-bc48-a051148e3d03 | 0.89                |
| 1 | 2022-03-01 09:00:00 | 01f3cdd9-8e9e-4dff-9b5c-69698a0388d0 | 0.14                |
| 2 | 2022-03-01 09:00:00 | 01ff0803-ae73-4234-971d-5713c97b7f4b | 0.67                |
| 3 | 2022-03-01 09:00:00 | 0363eb21-8c74-47e1-a216-c37e565e5ceb | 0.82                |
| 4 | 2022-03-01 09:00:00 | 03f0b20e-3b5b-444f-bc39-cdfa2523d4bc | 0.05                |

```
In [34]: sensor_storage_temperature_agg = sensor_storage_temperature.groupby(['timestamp']).agg({'temperature': 'mean'})  
sensor_storage_temperature_agg.head()
```

```
Out[34]:
```

|   | timestamp           | temperature |
|---|---------------------|-------------|
| 0 | 2022-03-01 09:00:00 | -0.028850   |
| 1 | 2022-03-01 10:00:00 | 1.284314    |
| 2 | 2022-03-01 11:00:00 | -0.560000   |
| 3 | 2022-03-01 12:00:00 | -0.537721   |
| 4 | 2022-03-01 13:00:00 | -0.188734   |

This gives us the average temperature of the storage facility where the produce is stored in the warehouse by unique hours during the week. Now, we are ready to merge our data. We will use the `stock_agg` table as our base table, and we will merge our other 2 tables onto this.

```
In [35]: merged_df = sensor_stock_levels_agg.merge(sales_agg, on=['timestamp', 'product_id'], how='left')  
merged_df.head()
```

```
Out[35]:
```

|   | timestamp           | product_id                           | estimated_stock_pct | quantity |
|---|---------------------|--------------------------------------|---------------------|----------|
| 0 | 2022-03-01 09:00:00 | 00e120bb-89d6-4df5-bc48-a051148e3d03 | 0.89                | 3.0      |
| 1 | 2022-03-01 09:00:00 | 01f3cdd9-8e9e-4dff-9b5c-69698a0388d0 | 0.14                | 3.0      |
| 2 | 2022-03-01 09:00:00 | 01ff0803-ae73-4234-971d-5713c97b7f4b | 0.67                | NaN      |
| 3 | 2022-03-01 09:00:00 | 0363eb21-8c74-47e1-a216-c37e565e5ceb | 0.82                | NaN      |
| 4 | 2022-03-01 09:00:00 | 03f0b20e-3b5b-444f-bc39-cdfa2523d4bc | 0.05                | NaN      |

```
In [36]: merged_df = merged_df.merge(sensor_storage_temperature_agg, on=['timestamp'], how='left')  
merged_df.head()
```

```
Out[36]:
```

|   | timestamp           | product_id                           | estimated_stock_pct | quantity | temperature |
|---|---------------------|--------------------------------------|---------------------|----------|-------------|
| 0 | 2022-03-01 09:00:00 | 00e120bb-89d6-4df5-bc48-a051148e3d03 | 0.89                | 3.0      | -0.02885    |
| 1 | 2022-03-01 09:00:00 | 01f3cdd9-8e9e-4dff-9b5c-69698a0388d0 | 0.14                | 3.0      | -0.02885    |
| 2 | 2022-03-01 09:00:00 | 01ff0803-ae73-4234-971d-5713c97b7f4b | 0.67                | NaN      | -0.02885    |
| 3 | 2022-03-01 09:00:00 | 0363eb21-8c74-47e1-a216-c37e565e5ceb | 0.82                | NaN      | -0.02885    |
| 4 | 2022-03-01 09:00:00 | 03f0b20e-3b5b-444f-bc39-cdfa2523d4bc | 0.05                | NaN      | -0.02885    |

```
In [37]: merged_df.isna().sum()
```

```
Out[37]:
```

|                     |       |
|---------------------|-------|
| timestamp           | 0     |
| product_id          | 0     |
| estimated_stock_pct | 0     |
| quantity            | 7778  |
| temperature         | 0     |
| dtype:              | int64 |

We can see from the `.isna` method that we have some null values. These need to be treated before we can build a predictive model. The column that features some null values is `quantity`. We can assume that if there is a null value for this column, it represents that there were 0 sales of this product within this hour. So, let's fill this column's null values with 0, however, we should verify this with the client, in order to make sure we're not making any assumptions by filling these null values with 0.

```
In [38]: merged_df['quantity'] = merged_df['quantity'].fillna(0)
merged_df.isna().sum()
```

```
Out[38]:
```

|                     |       |
|---------------------|-------|
| timestamp           | 0     |
| product_id          | 0     |
| estimated_stock_pct | 0     |
| quantity            | 0     |
| temperature         | 0     |
| dtype:              | int64 |

We now add these other columns, `category` and `unit_price`.

```
In [39]: product_categories = sales_data[['product_id', 'category']]
product_categories = product_categories.drop_duplicates()

product_price = sales_data[['product_id', 'unit_price']]
product_price = product_price.drop_duplicates()
product_price
```

```
Out[39]:
```

|      | product_id                            | unit_price |
|------|---------------------------------------|------------|
| 0    | 3bc6c1ea-0198-46de-9ffd-514ae3338713  | 3.99       |
| 1    | ad81b46c-bf38-41cf-9b54-5fe7f5eba93e  | 3.99       |
| 2    | 7c55cbd4-f306-4c04-a030-628cbe7867c1  | 0.19       |
| 3    | 80da8348-1707-403f-8be7-9e6deecccc883 | 0.19       |
| 4    | 7f5e86e6-f06f-45f6-bf44-27b095c9ad1d  | 4.49       |
| ...  | ...                                   | ...        |
| 7569 | d6ccd088-11be-4c25-aa1f-ea87c01a04db  | 14.99      |
| 7570 | 20a9bd7b-daff-4b8b-bdc1-2e8f9a0277fa  | 13.49      |
| 7572 | a9325c1a-2715-41df-b7f4-3078fa5ecd97  | 14.19      |
| 7576 | 0e4c10f4-77bc-4c67-86b2-b4da5ded19bf  | 16.99      |
| 7579 | bc6187a9-d508-482b-9ca6-590d1cc7524f  | 14.19      |

300 rows × 2 columns

```
In [40]: merged_df = merged_df.merge(product_categories, on="product_id", how="left")
merged_df = merged_df.merge(product_price, on="product_id", how="left")
merged_df.head()
```

```
Out[40]:
```

|   | timestamp           | product_id                           | estimated_stock_pct | quantity | temperature | category      | unit_price |
|---|---------------------|--------------------------------------|---------------------|----------|-------------|---------------|------------|
| 0 | 2022-03-01 09:00:00 | 00e120bb-89d6-4df5-bc48-a051148e3d03 | 0.89                | 3.0      | -0.02885    | kitchen       | 11.19      |
| 1 | 2022-03-01 09:00:00 | 01f3cdd9-8e9e-4dff-9b5c-69698a0388d0 | 0.14                | 3.0      | -0.02885    | vegetables    | 1.49       |
| 2 | 2022-03-01 09:00:00 | 01ff0803-ae73-4234-971d-5713c97b7f4b | 0.67                | 0.0      | -0.02885    | baby products | 14.19      |
| 3 | 2022-03-01 09:00:00 | 0363eb21-8c74-47e1-a216-c37e565e5ceb | 0.82                | 0.0      | -0.02885    | beverages     | 20.19      |
| 4 | 2022-03-01 09:00:00 | 03f0b20e-3b5b-444f-bc39-cdfa2523d4bc | 0.05                | 0.0      | -0.02885    | pets          | 8.19       |

## Section 9 - Feature engineering

We have our cleaned and merged data. Now we must transform this data so that the columns are in a suitable format for a machine learning model. In other terms, every column must be numeric.

Let's first engineer the `timestamp` column. In its current form, it is not very useful for a machine learning model. Since it's a datetime datatype, we can explode this column into day of week, day of month and hour to name a few.

```
In [41]: merged_df['timestamp_day_of_month'] = merged_df['timestamp'].dt.day
merged_df['timestamp_day_of_week'] = merged_df['timestamp'].dt.dayofweek
merged_df['timestamp_hour'] = merged_df['timestamp'].dt.hour
merged_df.drop(columns=["product_id", "timestamp"], inplace=True)
merged_df.head()
```

```
Out[41]:
```

|   | estimated_stock_pct | quantity | temperature | category      | unit_price | timestamp_day_of_month | timestamp_day_of_week | timestamp_hour |
|---|---------------------|----------|-------------|---------------|------------|------------------------|-----------------------|----------------|
| 0 | 0.89                | 3.0      | -0.02885    | kitchen       | 11.19      | 1                      | 1                     | 9              |
| 1 | 0.14                | 3.0      | -0.02885    | vegetables    | 1.49       | 1                      | 1                     | 9              |
| 2 | 0.67                | 0.0      | -0.02885    | baby products | 14.19      | 1                      | 1                     | 9              |
| 3 | 0.82                | 0.0      | -0.02885    | beverages     | 20.19      | 1                      | 1                     | 9              |
| 4 | 0.05                | 0.0      | -0.02885    | pets          | 8.19       | 1                      | 1                     | 9              |

The next column that we can engineer is the `category` column. In its current form it is categorical. We can convert it into numeric by creating dummy variables from this categorical column.

```
In [42]: merged_df = pd.get_dummies(merged_df, columns=['category'])
merged_df
```

```
Out[42]:
```

|       | estimated_stock_pct | quantity | temperature | unit_price | timestamp_day_of_month | timestamp_day_of_week | timestamp_hour | category_b... |
|-------|---------------------|----------|-------------|------------|------------------------|-----------------------|----------------|---------------|
| 0     | 0.89                | 3.0      | -0.028850   | 11.19      | 1                      | 1                     | 9              |               |
| 1     | 0.14                | 3.0      | -0.028850   | 1.49       | 1                      | 1                     | 9              |               |
| 2     | 0.67                | 0.0      | -0.028850   | 14.19      | 1                      | 1                     | 9              |               |
| 3     | 0.82                | 0.0      | -0.028850   | 20.19      | 1                      | 1                     | 9              |               |
| 4     | 0.05                | 0.0      | -0.028850   | 8.19       | 1                      | 1                     | 9              |               |
| ...   | ...                 | ...      | ...         | ...        | ...                    | ...                   | ...            | ...           |
| 10840 | 0.50                | 4.0      | -0.165077   | 4.99       | 7                      | 0                     | 19             |               |
| 10841 | 0.26                | 0.0      | -0.165077   | 19.99      | 7                      | 0                     | 19             |               |
| 10842 | 0.78                | 3.0      | -0.165077   | 6.99       | 7                      | 0                     | 19             |               |
| 10843 | 0.92                | 3.0      | -0.165077   | 14.99      | 7                      | 0                     | 19             |               |
| 10844 | 0.01                | 2.0      | -0.165077   | 5.19       | 7                      | 0                     | 19             |               |

10845 rows × 29 columns

## Section 10 - Modelling

Now it is time to train a machine learning model. We used a supervised machine learning model using `estimated_stock_pct` as the target variable, since the problem statement was focused on being able to predict the stock levels of products on an hourly basis.

Whilst training the machine learning model, we will use cross-validation, which is a technique where we hold back a portion of the dataset for testing in order to compute how well the trained machine learning model is able to predict the target variable.

Finally, to ensure that the trained machine learning model is able to perform robustly, we will want to test it several times on random samples of data, not just once. Hence, we will use a K-fold strategy to train the machine learning model on K (K is an integer to be decided) random samples of the data. First, let's create our target variable y and independent variables X

```
In [43]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error
from sklearn.preprocessing import StandardScaler
```

```
In [44]: features = merged_df.drop(columns=['estimated_stock_pct'])
target = merged_df['estimated_stock_pct']
print(features.shape)
print(target.shape)
```

```
(10845, 28)
(10845,)
```

For this exercise, we are going to use a `RandomForestRegressor` model, which is an instance of a `Random Forest`. These are powerful

For this exercise, we are going to use a Random Forest regressor model, which is an instance of a Random Forest. These are powerful tree based ensemble algorithms and are particularly good because their results are very interpretable.

We are using a regression algorithm here because we are predicting a continuous numeric variable, that is, estimated\_stock\_pct. A classification algorithm would be suitable for scenarios where you're predicted a binary outcome, e.g. True/False

```
In [45]: K = 10
        split = 0.75
```

```
In [46]: accuracy = []

for fold in range(0, K):

    # Instantiate algorithm
    model = RandomForestRegressor()
    scaler = StandardScaler()

    # Create training and test samples
    features_train, features_test, target_train, target_test = train_test_split(features, target, train_size=split)

    # Scale X data, we scale the data because it helps the algorithm to converge
    # and helps the algorithm to not be greedy with large values
    scaler.fit(features_train)
    X_train = scaler.transform(features_train)
    X_test = scaler.transform(features_test)

    # Train model
    trained_model = model.fit(features_train, target_train)

    # Generate predictions on test sample
    y_pred = trained_model.predict(features_test)

    # Compute accuracy, using mean absolute error
    mae = mean_absolute_error(target_test, y_pred)
    accuracy.append(mae)
    print(f"Fold {fold + 1}: MAE = {mae:.3f}")

print(f"Average MAE: {(sum(accuracy) / len(accuracy)):.2f}")
```


```
Fold 1: MAE = 0.237
Fold 2: MAE = 0.236
Fold 3: MAE = 0.236
Fold 4: MAE = 0.236
Fold 5: MAE = 0.236
Fold 6: MAE = 0.236
Fold 7: MAE = 0.236
Fold 8: MAE = 0.237
Fold 9: MAE = 0.237
Fold 10: MAE = 0.236
Average MAE: 0.24
```

```
In [47]: features = [i.split()[0] for i in features.columns]
        importances = model.feature_importances_

        indices = np.argsort(importances)

        fig, ax = plt.subplots(figsize=(10, 20))
        plt.title('Feature Importances')
        plt.barh(range(len(indices)), importances[indices], color='b', align='center')
        plt.yticks(range(len(indices)), [features[i] for i in indices])
        plt.xlabel('Relative Importance')
        plt.show()
```





This feature importance visualisation tells us:

1. The product categories were not that important
2. The unit price and temperature were important in predicting stock
3. The hour of day was also important for predicting stock